

SemEval 2022: Patronizing and Condescending Language Detection

Raluca Gînga **Bogdan Dobre** **Radu Sielecki** **Tudor Dumitraşcu**
University of Bucharest University of Bucharest University of Bucharest University of Bucharest

Abstract

This report is part of the SemEval 2022 Workshop, Task 4 - Patronizing and Condescending Language Detection, trying to find out Patronizing and Condescending Language in any form of text. There were used many methods, varying from simple Machine Learning algorithms applied on bag of words embeddings until Bert Embeddings and using Neural Networks in order to solve both the binary classification and multi-label classification as well.

1 Introduction

The Patronizing and Condescending Language Detection Task is based on the paper Don't Patronize Me! An annotated Dataset with Patronizing and Condescending Language Towards Vulnerable Communities (Pérez-Almendros et al., 2020).

The aim of this task is to identify PCL, and to categorize the linguistic techniques used to express it, specifically when referring to communities identified as being vulnerable to unfair treatment in the media.

Participants were provided with sentences in context (paragraphs), extracted from news articles, in which one or several predefined vulnerable communities are mentioned. The challenge is divided into two subtasks.

1. Subtask 1: Binary classification. Given a paragraph, a system must predict whether or not it contains any form of PCL.
2. Subtask 2: Given a paragraph, a system must identify which PCL categories express the condescension. The PCL taxonomy has been defined based on previous works on PCL. There are considered the following categories:

- Unbalanced power relations.
- Shallow solution.

- Presupposition.
- Authority voice.
- Metaphor.
- Compassion.
- The poorer, the merrier.

2 Background

The dataset used for this SemEval 2022 task was Don't Patronize Me! dataset, which contains a suite of sentences that mention some vulnerable communities and published in media in a lot of English speaking countries. The paragraphs were manually annotated to show 1) whether the text contains any kind of PCL, and 2) if it contains PCL, what linguistic techniques (categories) are used to express the condescension.

The dataset for subtask 1 (binary classification) contained a number of 10.636 paragraphs and 2.792 instances were used for the categories classification subtask.

In Figure 1, it can be seen that for the first subtask, there are almost 1000 of texts that contain PCL. That means we're dealing with imbalanced data that we need to solve it. In the next 3 figures (2, 3, 4), it could be noticed the distribution of the most common words both in the full dataset, but in texts that contain/don't contain PCL as well.

For tasks 2, the paragraphs from task 1 are split according to the type of PCL speech category into sentences, resulting in 950 samples.

3 System Overview

1. Subtask 1 (Binary Classification)

Because the dataset was very imbalanced, we tried different approaches in order to make it balanced:




Figure 1: Classes Distribution for Binary Classification problem (Subtask 1)




Figure 3: Most common words classified into PCL

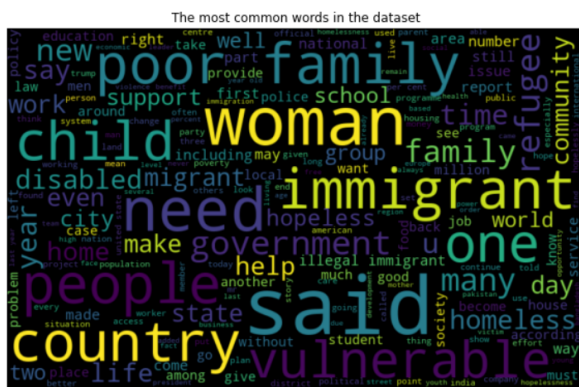


Figure 2: Most common words in the dataset (Subtask 1)



Figure 4: Most common words of texts that are not PCL

- Adding a class weight to the models used. In this approach, we computed a metric in which we obtained a class weight according to the imbalance of the dataset. Through this method, we gave some different weights to both the majority and minority classes. This whole process had the purpose to penalize the misclassification made by the minority class by setting a higher class weight and at the same time, reducing the weight for the majority class.
- Using oversampling methods and special ensemble techniques. In this approach, there were used methods like SMOTE (Synthetic Minority Over-sampling Technique), Adasyn (Adaptive Synthetic), SVM-SMOTE and self paced ensemble that performs strictly balanced under-sampling in each iteration, being very efficient computationally.
- Augmenting the data. Because we notice

so little data for label 1, we decided to collect hate speech datasets and add the positive texts into our dataset in order to balance the classes frequency, obtaining a total of 6372 from 795 initial texts with label 1. We'll notice in the results section that this collection and generation of new dataset didn't provide good results.

The dataset was a little bit preprocessed and split into two preprocessed types: lemmatized cleaned dataset and stemmed cleaned dataset. These two datasets were generated in order to make some comparison between those two techniques and to see which provided the best results.

As feature extraction techniques, there were used techniques like Bag of Words, Tokenizer, Word2Vec and, finally, BertTokenizer which provided the best results in the end.

As models, there were used Neural Networks with 3 dense layers, Long Short Term Memory (LSTM) with 64 and 128 neurons with

dropout as well, basic Machine Learning algorithms like Logistic Regression, Random Forest, Support Vector Machines as XGBoost. In the end, we decided to try BERT embeddings and a classification BERT model, Bert-ForSequenceClassification, that contains a single linear classification layer on top and that provided the best results after all of the other approaches.

Another approach, called "Text shards" made use of the subtask related to multiclass classification as well. For an average text that contains PCL, only some small pieces of them are actually PCL and the rest of the text are not. The assumption is that this confuses the model, because a combination of pcl and non pcl is labeled as PCL. To address this, the following approach is used:

- negative examples are left as they are
- each positive example is replaced with the actual pieces of PCL inside it that we can get from the categories file
- the positive examples obtained this way are added with the negative examples to obtain a training dataset
- all the sentences are cleaned of characters that are not letters and the words in each sentence are lemmatized
- a Tensorflow Hub pretrained model called Universal Sentence Encoder is trained on it
- for each text that we want to predict, we first use the model on the whole text to get an initial label
- a window (of the size of the average length of a cleaned PCL fragment * 2) is slid through the text and the model is used to predict that particular substring. If it is labeled as PCL, then we consider the whole text as PCL.

2. Subtask 2 (Multi-label Classification)

Considering the fact that the vocabulary of the English language is large, we have tried to leverage the power of pretrained language models. Therefore we have chosen 3 bert-based models which were pretrained for hate speech detection and sentiment analysis. The BERT models also provided a tokenizer which split

the sentences into tokens and appended the required tokens.

- BERT (Devlin et al., 2018) Uncased
- BERT Multilingual Uncased
- BERT HateXplain (Mathew et al., 2020): This model was trained to classify text as Hate speech, Offensive or Normal. It was trained on Gab, Twitter and Humain Rationale;
- Distil BERT (?): This model is a version of Distilled bert finetuned on the Twitter dataset;
- Distil BERT Multilingual Cased (Sanh et al., 2019)
- Distill RoBERTa (?): This model is a version of Distilled RoBERTa finetuned on the Twitter dataset;

In the paper describing the dataset (Pérez-Almendros et al., 2020), the authors group the categories into 3 General categories.

- (a) The saviour: Unbalanced power relations and Shallow relations
- (b) The Expert: Presupposition and Authority voice
- (c) The Poet: Compassion, Metaphor and The poorer the merrier

From this idea, we tried to train the models to predict those 3 categories, and save the hidden features to a fixed latent space. Then these learned features can be used when training the model to predict the required 7 sub-classes.

Along with those bert-based model, we also tried to implement models based on Word2Vec (Goldberg and Levy, 2014) trained on "Google News" and Machine Learning algorithms based on TF-IDF and BOW:

- LSTM Word2Vec Embeddings (Staudemeyer and Morris, 2019)
- biLSTM Word2Vec Embeddings (Huang et al., 2015)
- RNN Word2Vec Embeddings (Sherstinsky, 2020)
- SVM TF-IDF
- RandomForest TF-IDF

We also dabbled with the thought of training our own Word2Vec, in order to create a model

specialized on hate speech. However we decided against this idea, due to the lack of usable datasets and the computational resources required for this task.

4 Results

1. Subtask 1 (Binary Classification)

(a) Deep Learning / Machine Learning for Imbalanced and Oversampled dataset

Approach	Lemmatized (F1_Score)			
	Simple	SPE	SMOTE	SVMMSMOTE
Neural Networks	0.27	-	0.2823	0.3187
Logistic Regression	0.34	-	0.35	0.35
Random Forest	0.067	0.31	0.19	0.16
Support Vector Machines	0.27	-	0.10	0.14
XGBoost	0.15	-	0.23	0.24

Approach	Stemmed (F1_Score)			
	Simple	SPE	SMOTE	SVMMSMOTE
Neural Networks	0.2698	-	0.289	0.3166
Logistic Regression	0.35	-	0.38	0.37
Random Forest	0.038	0.31	0.21	0.13
Support Vector Machines	0.27	-	0.14	0.20
XGBoost	0.17	-	0.23	0.24

(b) Tokenization of the lemmatized and stemmed dataset / Word2Vec + LSTM neural network

Approach	Lemmatized (F1_Score)		Stemmed (F1_Score)	
	Simple	Word2Vec	Simple	Word2Vec
LSTM (64 neurons)	0.2693	0.2109	0.3213	0.2093
LSTM (128 neurons)	0.2317	0.2308	0.2789	0.2412

(c) Data augmentation

Augmented Data (F1_Score)	
NNs with 3 layers	0.2155
Logistic Regression	0.23
U.S.E. + 2 dense layers	0.2316

(d) BERT Transformers + BertForSequence-Classification

Unprocessed data (F1_Score)	
Bert Tokenizer + Classifier	0.5074

2. Subtask 2 (Multiclass Classification)

In Table 1, it can be seen that the models can easily two of the categories, while the rest of them are not learned. In Table 2 the general class approach is detailed, and even though some classes were learned, when it came to transfer the knowledge, information was lost.

5 Conclusion

In the current project, it was solved the problem posed by SemEval 2022 Task 4: Patronizing and Condescending Language Detection. There were applied various methods, including the application of Word Embeddings (Bag of Words, Word2Vec, BERT), tokenization, oversampling/undersampling of the datasets.

In the binary classification problem, the approach that gave the best result on the validation

Model	Unb	Sha	Pre	Aut
Bert	0.82	0.0	0.0	0.0
DistilRoberta	0.83	0.0	0.0	0.0
DistilBert	0.82	0.0	0.0	0.0

Model	Met	Com	Mer	F1-Mean
Bert	0.0	0.0	0.64	0.21
DistilRoberta	0.0	0.0	0.59	0.
DistilBert	0.66	0.08	0.0	0.34

Table 1: F1 Scores across the categories

Model	Expert		
	Class	Aut	Pre
DistilRoberta	0.54	0.0	0.0
DistilBertMLC	0.36	0.0	0.0
DistilBertUE	0.42	0.0	0.40
Bert	0.44	0.0	0.0

	Saviour		
	Class	Sha	Unb
DistilRoberta	0.85	0.0	0.84
DistilBertMLC	0.86	0.0	0.83
DistilBertUE	0.75	0.0	0.81
Bert	0.85	0.0	0.84

	Poet			
	Class	Com	Merr	Met
DistilRoberta	0.69	0.11	0.0	0.65
DistilBertMLC	0.60	0.0	0.0	0.52
DistilBertUE	0.61	0.0	0.0	0.67
Bert	0.69	0.0	0.0	0.59

Table 2: F1 Scores with sub-classes approach

dataset was BERT transformers combined with Bert for Sequence Classification, obtaining 0.50 as f1_score.

In the multi classification multi label task, the number of labels proved to be a challenge. The results overall are low and the models are only able to learn only a few classes. The general class approach also proved to be inefficient.

References

- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. [BERT: pre-training of deep bidirectional transformers for language understanding](#). *CoRR*, abs/1810.04805.
- Yoav Goldberg and Omer Levy. 2014. [word2vec explained: deriving mikolov et al.'s negative-sampling word-embedding method](#).

- Zhiheng Huang, Wei Xu, and Kai Yu. 2015. [Bidirectional lstm-crf models for sequence tagging](#).
- Binny Mathew, Punyajoy Saha, Seid Muhie Yimam, Chris Biemann, Pawan Goyal, and Animesh Mukherjee. 2020. Hatexplain: A benchmark dataset for explainable hate speech detection. *arXiv preprint arXiv:2012.10289*.
- Carla Pérez-Almendros, Luis Espinosa-Anke, and Steven Schockaert. 2020. [Don't patronize me! an annotated dataset with patronizing and condescending language towards vulnerable communities](#).
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *ArXiv*, abs/1910.01108.
- Alex Sherstinsky. 2020. [Fundamentals of recurrent neural network \(rnn\) and long short-term memory \(lstm\) network](#). *Physica D: Nonlinear Phenomena*, 404:132306.
- Ralf C. Staudemeyer and Eric Rothstein Morris. 2019. [Understanding lstm – a tutorial into long short-term memory recurrent neural networks](#).