

SemEval 2022: Patronizing and Condescending Language Detection

Andreea

Bogdan

Radu

Tudor

Abstract

This report is part of the SemEval 2022 Workshop, Task 4 - Patronizing and Condescending Language Detection, trying to find out Patronizing and Condescending Language in any form of text. There were used many methods, varying from simple Machine Learning algorithms applied on bag of words embeddings until Bert Embeddings and using Neural Networks in order to solve both the binary classification and multi-label classification as well.

1 Introduction

The Patronizing and Condescending Language Detection Task is based on the paper Don't Patronize Me! An annotated Dataset with Patronizing and Condescending Language Towards Vulnerable Communities (Perez-Almendros et al., 2020).

The aim of this task is to identify PCL, and to categorize the linguistic techniques used to express it, specifically when referring to communities identified as being vulnerable to unfair treatment in the media.

Participants were provided with sentences in context (paragraphs), extracted from news articles, in which one or several predefined vulnerable communities are mentioned. The challenge is divided into two subtasks.

1. Subtask 1: Binary classification. Given a paragraph, a system must predict whether or not it contains any form of PCL.
2. Subtask 2: Given a paragraph, a system must identify which PCL categories express the condescension. The PCL taxonomy has been defined based on previous works on PCL. There are considered the following categories:

- Unbalanced power relations.
- Shallow solution.

- Presupposition.
- Authority voice.
- Metaphor.
- Compassion.
- The poorer, the merrier.

2 Background

The dataset used for this SemEval 2022 task was Don't Patronize Me! dataset, which contains a suite of sentences that mention some vulnerable communities and published in media in a lot of English speaking countries. The paragraphs were manually annotated to show 1) whether the text contains any kind of PCL, and 2) if it contains PCL, what linguistic techniques (categories) are used to express the condescension.

The dataset for subtask 1 (binary classification) contained a number of 10.636 paragraphs and 2.792 instances were used for the categories classification subtask.

In Figure 1, it can be seen that for the first subtask, there are almost 1000 of texts that contain PCL. That means we're dealing with imbalanced data that we need to solve it. In the next 3 figures (2, 3, 4), it could be noticed the distribution of the most common words both in the full dataset, but in texts that contain/don't contain PCL as well.

For task 2, the paragraphs from task 1 are split according to the type of PCL speech category into sentences, resulting in 950 samples.

3 System Overview

1. Subtask 1 (Binary Classification)

Because the dataset was very imbalanced, we tried different approaches in order to make it balanced:



Figure 1: Classes Distribution for Binary Classification problem (Subtask 1)



Figure 2: Most common words in the dataset (Subtask 1)

- Adding a class weight to the models used. In this approach, we computed a metric in which we obtained a class weight according to the imbalance of the dataset. Through this method, we gave some different weights to both the majority and minority classes. This whole process had the purpose to penalize the misclassification made by the minority class by setting a higher class weight and at the same time, reducing the weight for the majority class.
- Using oversampling methods and special ensemble techniques. In this approach, there were used methods like SMOTE (Synthetic Minority Over-sampling Technique), Adasyn (Adaptive Synthetic), SVM-SMOTE and self paced ensemble that performs strictly balanced under-sampling in each iteration, being very efficient computationally.
- Augmenting the data. Because we notice so little data for label 1, we decided to collect hate speech datasets and add the positive texts into our dataset in order to balance the classes frequency, obtaining a total of 6372 from 795 initial texts with label 1. We'll notice in the results section that this collection and generation of new dataset didn't provide good results.

The dataset was a little bit preprocessed and split into two preprocessed types: lemmatized cleaned dataset and stemmed cleaned dataset. These two datasets were generated in order to make some comparison between those two techniques and to see which provided the best results.

As feature extraction techniques, there were used techniques like Bag of Words, Tokenizer, Word2Vec and, finally, BertTokenizer which provided the best results in the end.

As models, there were used Neural Networks with 3 dense layers, Long Short Term Memory (LSTM) with 64 and 128 neurons with dropout as well, basic Machine Learning algorithms like Logistic Regression, Random Forest, Support Vector Machines as XGBoost. In the end, we decided to try BERT embeddings and a classification BERT model, BertForSequenceClassification, that contains a single linear classification layer on top and that provided the best results after all of the other approaches.

Another approach, called "Text shards" made use of the subtask related to multiclass classification as well. For an average text that contains PCL, only some small pieces of them are actually PCL and the rest of the text are not. The assumption is that this confuses the model, because a combination of pcl and




Figure 3: Most common words classified into PCL



Figure 4: Most common words of texts that are not PCL

non pcl is labeled as PCL. To address this, the following approach is used:

- negative examples are left as they are
- each positive example is replaced with the actual pieces of PCL inside it that we can get from the categories file
- the positive examples obtained this way are added with the negative examples to obtain a training dataset
- all the sentences are cleaned of characters that are not letters and the words in each sentence are lemmatized
- a Tensorflow Hub pretrained model called Universal Sentence Encoder is trained on it
- for each text that we want to predict, we first use the model on the whole text to get an initial label
- a window (of the size of the average length of a cleaned PCL fragment * 2) is slid through the text and the model is used to predict that particular substring. If it is labeled as PCL, then we consider the whole text as PCL.

2. Subtask 2 (Multiclass Classification)

4 Results

Main quantitative findings: How well did your system perform at the task according to official metrics? How does it rank in the competition?

Quantitative analysis: Ablations or other comparisons of different design decisions to better understand what works best. Indicate which data split is used for the analyses (e.g. in table captions). If you modify your system subsequent to the official submission, clearly indicate which results are from the modified system.

Error analysis: Look at some of your system predictions to get a feel for the kinds of mistakes it makes. If appropriate to the task, consider including a confusion matrix or other analysis of error you may need to manually tag a small sample for this.

5 Conclusion

5.1 References

A Example Appendix

This is an appendix.